

Liability, Verification, and the Market for Expertise in the Presence of AI

Internal research seminar — Facultad de Ciencias Empresariales, Universidad Austral

Ricardo A. Pasquini^{a,b}

September 2026

^a Facultad de Ciencias Empresariales, Universidad Austral

^b Escuela de Gobierno, Universidad Torcuato Di Tella

`rpasquini@austral.edu.ar`

Motivation

The Model

Baseline Results

Capital-Constrained Bonding

Discussion, Extensions, and Conclusion

Motivation

The question

- Should a firm hire a consultant to run an analysis, or run it itself with an AI model?
- Used to turn mostly on **capability**: could the firm produce a comparable result at all?
- Increasingly turns on something else: can the firm **tell**, once the work is done, whether it is any good?

Two classes of tasks

Verifiable

Webpage, image/video, small software feature, first-pass slide deck — principal can click through it, run it, compare it to a brief.

Not verifiable

Tax opinion on a novel transaction, financial-statement audit, medical diagnosis — judging the output requires much the same expertise that produced it.

- The first class is drifting toward direct AI use. The second is where hiring survives — and where a problem is opening up.

Motivation, and what the model does

1. **Verification.** Unable to verify quality, the principal relies on a correlate (a polished deliverable, a fluent report). AI makes these cheap for anyone: correlates deteriorate, verification gets harder, and markets unravel when quality cannot be separated (Akerlof).
2. **A bond.** A genuine expert can propose a bond, forfeited on a provable failure: a credible signal even when the principal cannot inspect the work.
3. **But** enforcement relies on verification too: a court's ability to establish that an audit or diagnosis was wrong is itself a verification problem, often harder than the principal's.

What the model does

It does not model the deterioration of signals. It takes verification as given and analyzes how signaling works when verification is weak: can a bond fix the signaling problem?

What is interesting

Easy to verify: the principal does it himself. Hard: he pays for a better chance of a good outcome.

Verification is given and does not change. A liability bond can solve the signaling problem, under some conditions.

What this talk does

1. Sets out a two-stage model: a make-or-buy decision (DIY vs. delegate) nested around a liability-bond signaling game. It formalizes the tension between bare AI use and a contracting professional who uses AI; a contracting firm that bundles liability is an extension, not solved here.
2. Solves the **baseline** version, without a limit on the genuine expert's capital: separation occurs only in an intermediate **band** of verifiability.
3. Adds a **capital constraint** on bonding: it raises the band's lower edge.

Related literature

- **Signaling.** Spence (1973), Riley (1979); Cho and Kreps (1987) intuitive criterion — applied here to a liability bond rather than to education.
- **Judgment-proof problem.** Shavell (1986) — applied here to the credibility of a *signal*, not to ex post compensation.
- **Credence goods / unraveling.** Akerlof (1970); Darby and Karni (1973).
- **AI + liability, different friction.** Gans (2026) — staged access/release timing for dual-use AI. Bryan and Gans (2026) — training data composition given anticipated human use.

This paper's contribution

How does verification operate in the delegation of tasks, and does liability survive as a signaling device once the technology that created the need for it also erodes the court's ability to enforce it?

The Model

- Risk-neutral principal faces a task worth S if delivered **Good**, $S - L$ if delivered **Bad**. $L > 0$ is the stake destroyed by a Bad outcome.

Two primitives, drawn by nature before either party acts

- **Verification level** $v \in [0, 1]$ — probability that inspection, if performed, reveals the true state. *Same primitive* whether inspection is run by the principal (DIY) or, indirectly, by a court enforcing a bond (delegation).
- **Intermediary's type** $\theta \in \{H, L\}$ — H a genuine tacit-knowledge expert, L a *wrapper* reselling vanilla AI access. Private to the intermediary; common prior $\mu = \Pr(H)$.

0. Nature draws v (common knowledge) and θ (private, prior μ)



1. Principal chooses **DIY** (succeeds w.p. p_0) or enters the **delegation market**




2. Intermediary posts bond $B(\theta) \geq 0$ (capacity-bounded); principal forms belief $\mu(B)$, chooses **Engage** at price t or **Walk away**



3. Outcome Good w.p. $p_\theta \equiv 1 - f_\theta$; if Bad, verified w.p. $e(v) - B$ forfeited; else loss uncompensated

Two ways to get the task done: outcomes and payoffs


DIY (principal operates the AI)



$$U_{\text{DIY}}(v) = S - c_0 - (1 - v)f_0L - vf_0\kappa$$


Two ways to get the task done: outcomes and payoffs

DIY (principal operates the AI)



$$U_{\text{DIY}}(v) = S - c_0 - (1 - v)f_0L - vf_0\kappa$$

Delegation (type θ , bond B)



$$U(\theta, B) = S - t - f_\theta L + f_\theta v B$$

intermediary: $\pi(\theta, B) = t - f_\theta v B$

- $f_0 \equiv 1 - p_0$ — principal's own (non-expert) failure probability.
- c_0 — his cost of operating the AI directly.
- κ — cost of a successful retry when inspection catches a Bad draw (Assumption 3, next slides).

$$U_{\text{DIY}}(v) = S - (1 - v)f_0L - c_0 - vf_0\kappa$$

- Affine in v : $U_{\text{DIY}}(v) = S - f_0L - c_0 + v f_0(L - \kappa)$.
- In the leading case $L > \kappa$ (stake exceeds retry cost), U_{DIY} is **increasing** in v : better verification makes catch-and-retry more valuable.

Principal engages an intermediary of type θ who has posted bond B :

$$\pi(\theta, B) = t - f_{\theta} e(v) B \quad (\text{intermediary})$$

$$U(\theta, B) = S - t - f_{\theta} L + f_{\theta} e(v) B \quad (\text{principal})$$

- $\bar{U}$ — principal's outside option. Held fixed in §3 (baseline); from §3.3 on, $\bar{U} = U_{\text{DIY}}(v)$.
- Types are ordered: $f_L > f_H > 0$.

Notation, at a glance

$\theta \in \{H, L\}, \mu = \Pr(H)$	type, prior
$p_\theta, f_\theta \equiv 1 - p_\theta$	success / failure probability by type, $f_L > f_H$
t	price paid to the intermediary (exogenous, common)
B	bond posted by the intermediary
$e(v)$	$\Pr(\text{verified} \mid \text{Bad})$ — enforcement, $= v$ in the baseline
$\bar{w}$	capacity bound on the bond H can credibly post
S, L	task value if Good; stake destroyed if Bad
v	principal's own verification precision
$p_0, f_0 \equiv 1 - p_0$	DIY success / failure probability
c_0, κ	DIY operating cost; cost of a caught-and-redone retry
$\bar{U}$	principal's outside option

Assumptions 1–2

Assumption 1 (Fixed price)

The price t is exogenous and common to both intermediary types.

Type affects the probability of success, not the price of operating the AI — consistent with raw AI execution itself being commoditized.

Assumption 2 (Common knowledge of verifiability)

v is known to both parties before the bond is posted.

Without it, the principal's DIY outside option, which depends on v , would be *his* private information: the intermediary would be uncertain about the principal's participation threshold as well. Private information on both sides, instead of one.

Assumption 3

Assumption 3 (Mode-specific remedy)

A caught Bad draw is remedied differently depending on who operates the AI. Under DIY, the principal redoes the work himself at cost κ . Under delegation, the principal cannot redo the intermediary's work; his remedy is the bond.

- Retry is the principal's self-correction technology; the bond is delegation's remedy.
- Reflects what each party brings: the principal can redo *his own* work when he catches a failure; what an intermediary offers is a higher probability of getting it right the *first* time.
- This is the substantive reason the model builds two separate mechanisms (retry, bond) rather than one.

Assumption 4 (Capacity bound)

The genuine type cannot credibly post a bond larger than $\bar{w}$.

- The judgment-proof problem, applied to the credibility of the *signal itself*, not merely to ex post compensation.

Assumption 5, and the enforcement technology

Assumption 5 (Equilibrium selection)

We focus on (i) pooling without a bond, if it clears the principal's participation constraint; otherwise (ii) the least-cost separating equilibrium, if it exists. Where neither exists, separation fails, and we do not characterize the outcome.

- Both rules are *selections*. Pooling at a positive bond can coexist with separation for some priors; we set it aside, since it asks the genuine type to pay for a bond that does not distinguish it from the wrapper.

Enforcement technology

$e(v) = v$ — the most direct reading of “if it’s hard for the principal to verify, it’s hard for a court to verify too” as *one* primitive. A general affine form, $e(v) = e_0 + (1 - e_0)v$, changes only the capacity threshold of §4.

Baseline Results

3.1 — Set-up

Suppose momentarily the intermediary is a **single, known** type, succeeding with probability p_1 (compare the principal's own p_0).

- For a given verification level v : does the principal prefer to delegate at the posted price, or do it himself?
- Not yet about liability at all — there is nothing to signal when type is known.

Proposition 1 (make-or-buy)

Proposition 1

Suppose p_1 is known and $L > \kappa$. There is a unique threshold

$$v^* = \frac{\Delta p \cdot L - \Delta c}{f_0(L - \kappa)}, \quad \Delta p \equiv p_1 - p_0, \quad \Delta c \equiv t - c_0,$$

below which delegating strictly beats DIY, above which delegation cannot beat DIY.

Proposition 1 — derivation

No type uncertainty $\Rightarrow$ nothing to signal $\Rightarrow$ under Assumption 1, a bond is a pure transfer the intermediary dislikes, so $B = 0$:

$$U_{\text{DEL}} = S - t - f_1 L, \quad f_1 \equiv 1 - p_1 \quad (\text{flat in } v, \text{ no retry term by Assumption 3})$$

Proposition 1 — derivation

No type uncertainty $\Rightarrow$ nothing to signal $\Rightarrow$ under Assumption 1, a bond is a pure transfer the intermediary dislikes, so $B = 0$:

$$U_{\text{DEL}} = S - t - f_1 L, \quad f_1 \equiv 1 - p_1 \quad (\text{flat in } v, \text{ no retry term by Assumption 3})$$

DIY is affine and, in the leading case, increasing in v :

$$U_{\text{DIY}}(v) = S - f_0 L - c_0 + v f_0 (L - \kappa)$$

Proposition 1 — derivation

No type uncertainty $\Rightarrow$ nothing to signal $\Rightarrow$ under Assumption 1, a bond is a pure transfer the intermediary dislikes, so $B = 0$:

$$U_{\text{DEL}} = S - t - f_1 L, \quad f_1 \equiv 1 - p_1 \quad (\text{flat in } v, \text{ no retry term by Assumption 3})$$

DIY is affine and, in the leading case, increasing in v :

$$U_{\text{DIY}}(v) = S - f_0 L - c_0 + v f_0 (L - \kappa)$$

Their difference is affine with slope $-f_0(L - \kappa) < 0$ when $L > \kappa$:

$$U_{\text{DEL}} - U_{\text{DIY}}(v) = \Delta p \cdot L - \Delta c - v f_0 (L - \kappa)$$

Proposition 1 — derivation

No type uncertainty $\Rightarrow$ nothing to signal $\Rightarrow$ under Assumption 1, a bond is a pure transfer the intermediary dislikes, so $B = 0$:

$$U_{\text{DEL}} = S - t - f_1 L, \quad f_1 \equiv 1 - p_1 \quad (\text{flat in } v, \text{ no retry term by Assumption 3})$$

DIY is affine and, in the leading case, increasing in v :

$$U_{\text{DIY}}(v) = S - f_0 L - c_0 + v f_0 (L - \kappa)$$


Their difference is affine with slope $-f_0(L - \kappa) < 0$ when $L > \kappa$:

$$U_{\text{DEL}} - U_{\text{DIY}}(v) = \Delta p \cdot L - \Delta c - v f_0 (L - \kappa)$$

Exactly one sign change $\Rightarrow$ the threshold v^* above. $\square$

Proposition 1 — intuition

- High verifiability makes catch-and-retry a **substitute** for expertise: if a bad draw is cheaply caught and redone, you don't need someone who gets it right the first time.
- Low verifiability makes retry worthless — you can't tell Good from Bad — collapsing the comparison to a static “is the quality premium worth the price premium” calculation.
- That static comparison is exactly the terrain on which a credence-good unraveling problem can bite.



3.2 — The bond as a signal

Now: type uncertainty, fixed enforcement e , fixed outside option $\bar{U}$.

- Marginal cost of posting bond B to type θ : $\partial(f_\theta e B)/\partial B = f_\theta e$ — strictly lower for H , since $f_L > f_H$.
- That gap alone — not an assumed cost-of-effort asymmetry — is single crossing, and it's derived from one primitive: each type's own probability of a bad draw.

Proposition 2 (the bond)

Proposition 2

Suppose a known wrapper is not worth engaging even at a zero fee, $S - f_L L \leq \bar{U}$. The least-cost bond separating H from L is

$$B_H^* = \frac{t}{f_L e},$$

and H is always willing to post it.

- The condition $S - f_L L \leq \bar{U}$ is the *scope* of the result, not an assumption of the model: it says a wrapper that reveals itself is not worth engaging, so it earns zero. In the nested model it reads $v \geq v_W$ (§3.3).

Proposition 2 — derivation

A revealed wrapper is engaged at bond B iff $S - t - f_L L + f_L eB \geq \bar{U}$. At the smallest such bond it earns

$$t - f_L eB = S - f_L L - \bar{U} \leq 0$$

so an identified L 's payoff is 0. (This is where the scope condition is used.)

Proposition 2 — derivation

A revealed wrapper is engaged at bond B iff $S - t - f_L L + f_L e B \geq \bar{U}$. At the smallest such bond it earns

$$t - f_L e B = S - f_L L - \bar{U} \leq 0$$

so an identified L's payoff is 0. (This is where the scope condition is used.) Mimicking H's bond is therefore deterred exactly when

$$t - f_L e B \leq 0 \iff B \geq \frac{t}{f_L e}$$

Proposition 2 — derivation

A revealed wrapper is engaged at bond B iff $S - t - f_L L + f_L e B \geq \bar{U}$. At the smallest such bond it earns

$$t - f_L e B = S - f_L L - \bar{U} \leq 0$$

so an identified L's payoff is 0. (This is where the scope condition is used.) Mimicking H's bond is therefore deterred exactly when

$$t - f_L e B \leq 0 \iff B \geq \frac{t}{f_L e}$$

H wants the *smallest* such bond (bonding is costly): $B_H^* = t/(f_L e)$. Higher separating bonds exist too, but the standard refinement (Riley; Cho and Kreps) selects the least-cost one.

Proposition 2 — derivation

A revealed wrapper is engaged at bond B iff $S - t - f_L L + f_L eB \geq \bar{U}$. At the smallest such bond it earns

$$t - f_L eB = S - f_L L - \bar{U} \leq 0$$

so an identified L's payoff is 0. (This is where the scope condition is used.) Mimicking H's bond is therefore deterred exactly when

$$t - f_L eB \leq 0 \iff B \geq \frac{t}{f_L e}$$

H wants the *smallest* such bond (bonding is costly): $B_H^* = t/(f_L e)$. Higher separating bonds exist too, but the standard refinement (Riley; Cho and Kreps) selects the least-cost one. Check H's own participation:

$$t - f_H eB_H^* = t\left(1 - \frac{f_H}{f_L}\right) > 0 \quad \text{automatic, given } f_L > f_H. \quad \square$$

Proposition 3 (invariance to enforcement)

Proposition 3

At B_H^* , the principal's payoff does not depend on e :

$$\frac{\partial U(H, B_H^*)}{\partial e} = 0,$$

provided enforcement is costless to run.

Derivation — one substitution:

$$U(H, B_H^*) = S - t - f_H L + f_H e \cdot \frac{t}{f_L e} = S - t - f_H L + \frac{f_H t}{f_L}$$

e **cancels exactly** between the grossed-up bond and the recovery term. $\square$

- **Intuition:** worse enforcement inflates the nominal bond exactly enough to offset the lower chance of collecting it, so better courts make separation more *feasible*, not more valuable; feasibility is where the capacity bound of §4 enters.

Proposition 4 (pooling threshold)

Proposition 4

Pooling at $B = 0$ clears the principal's participation constraint iff $\mu \geq \mu^*$, where

$$\mu^* = \frac{f_L - \frac{S - t - \bar{U}}{L}}{f_L - f_H}.$$

Derivation. At $B = 0$, posterior = prior, so

$$U_{\text{pool}}(\mu) = S - t - [\mu f_H + (1 - \mu)f_L]L, \quad \text{increasing in } \mu \text{ since } f_H < f_L.$$

Set $U_{\text{pool}}(\mu) = \bar{U}$, solve for μ . $\square$

- If the prior is good enough, the principal is willing to engage even without separation.

Proposition 4 — intuition

- Above μ^* , H has no reason to signal: already engaged at $B = 0$ earning t ; any bond only lowers that.
- Follows from **Assumption 1** (fixed price): unlike Spence (1973), being identified as H is worth *nothing* beyond being engaged, because the price doesn't respond to belief.

3.3 — Nesting the two

Substitute $\bar{U} = U_{\text{DIY}}(v)$ and $e(v) = v$ into Propositions 2–4: the delegation market, screened by a bond, evaluated against an outside option that depends on the *same* verification parameter.

Define the principal's expected shortfall at $v = 0$, relative to a wrapper working for free:

$$\Delta \equiv Lf_0 + c_0 - Lf_L$$

Proposition 2's scope condition, $S - f_L L \leq U_{\text{DIY}}(v)$, becomes $v \geq v_W \equiv \Delta / [f_0(L - \kappa)]$.

- **Why v_W :** $S - f_L L$ is what a wrapper working for free would deliver. $U_{\text{DIY}}(v)$ rises with v , so it overtakes that at v_W . Below v_W , even a known wrapper is worth hiring for the principal, if it posts a large enough bond. So a wrapper that reveals itself still earns a positive payoff, and the logic of Prop. 2 (mimicking is deterred because a revealed wrapper earns nothing) fails.

Proposition 5 (the separating band)

Proposition 5

Suppose $L > \kappa$. Then

$$\mu^*(v) = \frac{f_L - f_0 + \frac{t-c_0}{L} + v \frac{f_0(L-\kappa)}{L}}{f_L - f_H}, \quad v_{\text{IR}} = \frac{f_0 L + c_0 - t - f_H L + f_H t / f_L}{f_0(L - \kappa)}$$

with $\mu^*(v)$ affine and increasing in v . Four regions:

- (a) *Pooling without a bond*: $\mu \geq \mu^*(v)$.
- (b) *Separating band*: $\mu < \mu^*(v)$, $v_W \leq v \leq v_{\text{IR}}$. H separates at B_H^* .
- (c) *No separation*: $\mu < \mu^*(v)$, $v < v_W$.
- (d) *DIY*: $v > v_{\text{IR}}$. Delegation cannot beat DIY.

Proposition 5 — derivation highlights

Substituting $\bar{U} = U_{\text{DIY}}(v)$ into Proposition 4's μ^* gives an affine map (algebra: direct substitution, slope $f_0(L - \kappa)/[L(f_L - f_H)] > 0$).

Proposition 5 — derivation highlights


Substituting $\bar{U} = U_{\text{DIY}}(v)$ into Proposition 4's μ^* gives an affine map (algebra: direct substitution, slope $f_0(L - \kappa)/[L(f_L - f_H)] > 0$). By Proposition 3, $U(H, B_H^*)$ is flat in v ; $U_{\text{DIY}}(v)$ is increasing. They cross once, at v_{IR} .

Proposition 5 — derivation highlights

Substituting $\bar{U} = U_{\text{DIY}}(v)$ into Proposition 4's μ^* gives an affine map (algebra: direct substitution, slope $f_0(L - \kappa)/[L(f_L - f_H)] > 0$). By Proposition 3, $U(H, B_H^*)$ is **flat** in v ; $U_{\text{DIY}}(v)$ is increasing. They cross once, at v_{IR} . Below v_W , an identified wrapper can obtain engagement and earn $S - f_L L - U_{\text{DIY}}(v) > 0$: the deterrence argument of Proposition 2 fails, so there is no separation.

Proposition 5 — intuition

- **Pooling without a bond** ($\mu \geq \mu^*(v)$): the prior is good enough for the principal to engage with no bond. $\mu^*(v)$ rises with v , because the DIY outside option improves.
- **Separating band**: H separates at B_H^* . Upper edge v_{IR} : Prop. 1 plus the bond.
- **No separation** ($v < v_W$): imitation is too good to screen out, even if H can post any bond.
- **DIY** ($v > v_{IR}$): catch-and-retry makes identifying an expert unnecessary.



Example of the last slides, $t = 1.2$

Reintroducing Assumption 4

- We now impose the capacity bound: H cannot post a bond larger than $\bar{w}$.
- The question: how does a limit on the bond change the separating band of Proposition 5?

Proposition 6 (band with capital constraint)

Proposition 6

Separation at B_H^* is feasible only if $B_H^* \leq \bar{w}$, i.e. $v \geq v_{\text{cap}} \equiv t/(f_L \bar{w})$. Under $e(v) = v$, the separating band becomes

$$\mu < \mu^*(v), \quad \max\{v_W, v_{\text{cap}}\} \leq v \leq v_{\text{IR}},$$

and is empty if $v_{\text{cap}} > v_{\text{IR}}$. For $\mu < \mu^*(v)$ and $v < \max\{v_W, v_{\text{cap}}\}$ there is no separation. The pooling and DIY regions of Proposition 5 are unchanged.

Proposition 6 — derivation

$$B_H^* = \frac{t}{f_L v} \leq \bar{w} \iff v \geq v_{\text{cap}} = \frac{t}{f_L \bar{w}} \quad \text{— a vertical line in } (v, \mu) \text{ space.}$$

Proposition 6 — derivation


$$B_H^* = \frac{t}{f_L v} \leq \bar{w} \iff v \geq v_{\text{cap}} = \frac{t}{f_L \bar{w}} \quad \text{— a vertical line in } (v, \mu) \text{ space.}$$

Combined with Proposition 5, H separates iff $\mu < \mu^*(v)$ and $\max\{v_W, v_{\text{cap}}\} \leq v \leq v_{\text{IR}}$.

Proposition 6 — intuition

- The less verifiable the task, the larger the bond needed to deter the wrapper ($B_H^* = t/(f_L v)$).
- The bond cannot exceed the expert's capital, so verifiability must be **high enough** for the bond to be payable.
- And it must be **low enough** that the principal does not simply do the task himself.
- Between the two, separation is possible: that is the band.

The (v, μ) map



Example: $f_H = 0.2$, $f_L = 0.5$, $f_0 = 0.8$, $L = 10$, $\kappa = 1$, $c_0 = 0.5$, $\bar{w} = 4$, $S = 20$, $t = 1.2$, giving $v_W = 0.486$, $v_{cap} = 0.600$, $v_{IR} = 0.803$. The hatched strip $[v_W, v_{cap})$ is where separation would occur if the expert could post any bond, and is lost to the capital constraint.

Discussion, Extensions, and Conclusion

Implications, a conjecture, and a caveat

- **Implications.** Better imitation (lower f_L) raises v_W and v_{cap} . More bonding capacity lowers v_{cap} .
- **Conjecture (not derived).** Where separation fails, the market cannot tell experts from imitators, and with a fixed price it does not reward expertise: it will eventually lose it.
- **Caveat.** Assumption 5 selects separation wherever pooling without a bond is not enough. Pooling clears the principal's IR only in expectation; a principal who cannot be compensated by a transfer, or is judged ex post, may reject it.

- **Costly bonds.** Real bonds carry costs $\propto$ face value (surety premiums, LC fees, collateral opportunity cost). Since $B_H^* \propto 1/e$, weaker enforcement makes separation more wasteful even when affordable — Prop. 3's invariance breaks.
- **Liability as one signal among several.** Reputation, credentials, institutional affiliation are alternatives with their own verification requirements, and AI erodes them through different channels; not compared here.
- **Risk and the ex-post cost of pooling.** The expected-value IR constraint suits the intermediary's repeated-trial view but is strong for a one-shot principal, especially where harm is not restorable by a transfer. Replacing it with a risk-averse or ex-post welfare criterion is left for future work.

Conclusion

Two mechanisms jointly determine whether a market for genuine expertise survives cheap, competent AI execution: **(i)** the principal's ability to catch and redo a bad result himself; **(ii)** a bond that lets genuine expertise be credibly signaled when he cannot.

The bond separates genuine expertise from imitation only in an **intermediate band of verifiability**.

- Above the band, the principal does the task himself. Below it, a wrapper cannot be screened out.
- A limit on the capital a genuine expert can post raises the band's lower edge further.
- Better imitation, which AI gives the wrapper, raises the edge; more bonding capacity lowers it wherever capital sets it.
- Where separation fails, the market cannot tell experts from imitators — we conjecture it will not keep its experts.

Thank you — questions and comments welcome.